

# Parallel Programming Tutorial - Pthread 1

M.Sc. Andreas Wilhelm  
 Technical University Munich  
 April 18, 2016



*TUM Uhrenturm*

# Organization

- Lecture could start at:
  - 4:00 PM
  - 4:15 PM
  - 4:30 PM
- Poll is available at <http://doodle.com/poll/cr6udw93bfvemwgu>
- Duration: as long as we need, up to 90 min

# Organization

- Assignments on parallel programming techniques
- Topics
  - Pthreads (Posix Threads)
  - OpenMP (Open Multi-Processing)
  - Dependency analysis
  - MPI (Message Passing Interface)
- Code examples are in C99 (no C++)
- My email address is: andreas.wilhelm(at)in.tum.de

# Assignments

- Submission of 80% of the assignments gives 0.3 bonus
- Submissions will be checked for:
  - plagiarism
  - correctness (output, threads, synchronization)
  - speedup
  - memory leaks
- Submission is done on website: <https://131.159.35.135/Submission>
  - requires your LRZ ID and your password
  - password is not stored and only used for authentication
  - download SSL-certificate here
- Assignment instructions are on the last slides
- Final exam will contain small programming tasks
- Solutions will be made public

# Assistance on Assignments

## After this week

- Given by: Helisa Dharmo and Amir Raoofy
- Email: helisa.dhamo(at)fti.edu.al / amir.raoofy(at)tum.de
- Room: 01.04.011
- Possible Date and Time:
  - Tuesday 10AM, 11AM, 12AM (60 mins.)
  - Thursday 10AM, 11AM, 12AM, 1PM, 2PM, 3PM (60 mins.)
  - Friday 10AM, 11AM, 12AM, 1PM, 2PM (60 mins.)
- Poll is available at <http://doodle.com/poll/asfpt8ge6zhzdvg9>


# Resources

- POSIX Threads Programming
- An Introduction to Parallel Programming, by Peter Pacheco
- Programming with Posix Threads, by David Butenhof
- The Linux Programming Interface, by Michael Kerrisk
- Patterns for Parallel Programming, by Timothy G. Mattson; Beverly A. Sanders; Berna L. Massingill

# Course Prerequisites

- Knowledge of C
  - memory management
  - pointers
  - global vs. static variables
- C books
  - (C89) The C Programming Language, Second Edition, by Brian W. Kernighan; Dennis M. Ritchie
  - (C99) C Primer Plus, Fifth Edition, by Stephen Prata
- Experience with Linux Command Line
- Resources
  - Book: The Linux Command Line
  - Basic video introduction: The Shell
- Knowing GCC
  - An Introduction to GCC, by Brian Gough

# 51th Anniversary of Moore's Law



Original data collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond and C. Batten  
Dotted line extrapolations by C. Moore



# Year 2005: The Free Lunch Is Over

- A Fundamental Turn Toward Concurrency in Software
- Software doesn't get (much) faster with the next microprocessor generation
- Software developers have to rewrite their applications to use multiple processors in order to speed them up
- Parallel Programming is hard
  - to write - complex APIs and needs more code than serial version
  - to do it correctly - it's easy to introduce new bugs
  - to debug, order of thread execution is undefined
  - to make it scalable - will your applications scale with more cores?
  - better qualified developers are necessary

# Posix Thread Programming

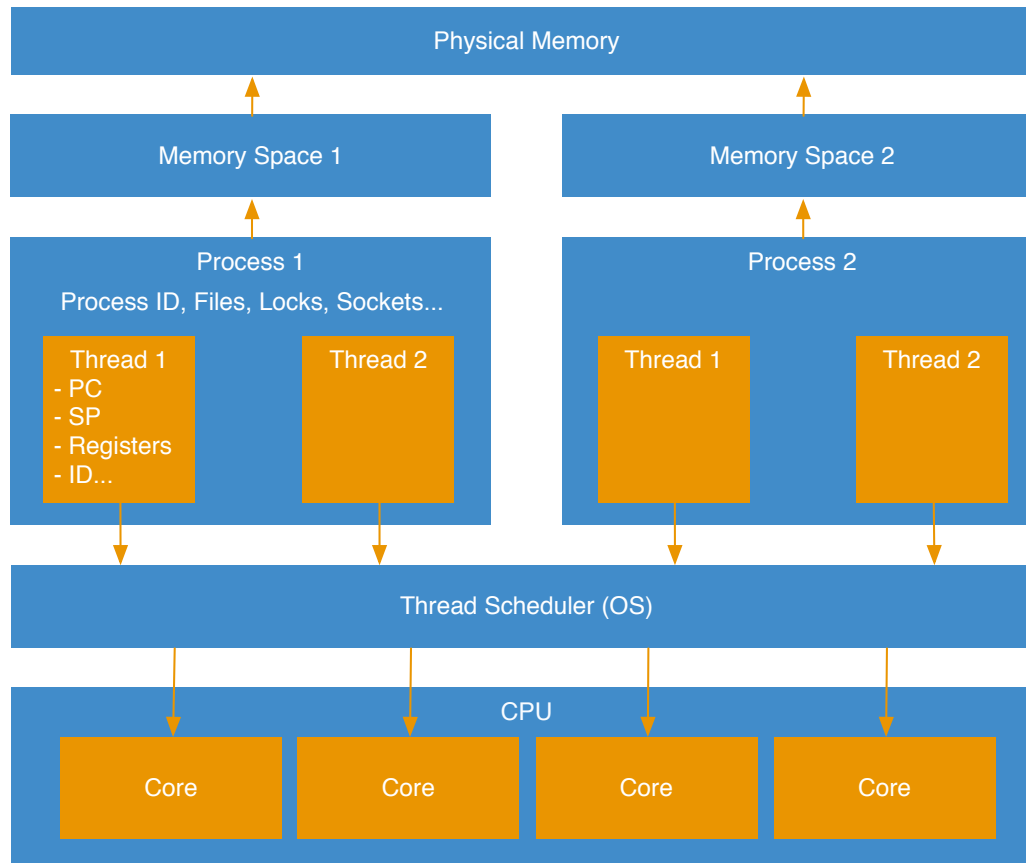
## Definition: Thread

A thread is an independent stream of instructions that can be scheduled to run as such by the operating system.

## POSIX Threads (Pthreads)

- Were defined in 1995 (IEEE Std 1003.1c-1995)
- Is an API that defines a set of types, functions and constants
- Is implemented with a `pthread.h` header and a thread library
- Functions can be categorized in four groups:
  - Thread management
  - Mutexes
  - Condition variables
  - Read/write locks and barriers

# Processes vs. Threads



# Why use Multithreading?

- **Performance gains**

Parallel processing by multiple processor cores

- **Increased application throughput**

Asynchronous system calls possible

- **Increased application responsiveness**

Application does not need to block operations

- **Replacing process-to-process communications**

Threads may communicate by shared-memory


- **Efficient use of system resources**

Lightweight context switches possible

- **The ability to create well-structured programs**

Some programs are inherently concurrent

# Concurrency vs. Parallelism



# Why are threads 'faster' than processes?

- Creating a new process with `fork()` has a big overhead: whole memory must be copied
  - Waste of memory space!
- Synchronization with processes usually involves system calls.

# Create Pthreads

```
1 int pthread_create(pthread_t *thread ,  
2                   const pthread_attr_t *attr ,  
3                   void *(*start_routine) (void *),  
4                   void *arg );
```

- pthread\_t \*thread,
  - Pointer to thread identifier.
- const pthread\_attr\_t \*attr
  - Optional pointer to pthread\_attr\_t to define behavior, if NULL defaults are used.
- void \*(\*start\_routine) (void \*),
  - Pointer to function prototype that is started. Function takes void pointer as argument and returns a void pointer.
- void \*arg
  - Pointer to the argument that is used for the executed function.

# Waiting for Pthread to finish

```
1 int pthread_join(pthread_t thread ,  
2                 void **retval);
```

- pthread\_t thread,
  - Pointer to thread identifier, for which this function is waiting.
- void \*\*retval
  - Optional pointer pointing to a void pointer. This can be used to return data of undefined size.



# Create Pthreads - Example

```
1 #include <stdio.h>
2 #include <pthread.h>
3
4 void* hello()
5 {
6     printf("Hello World from pthread!\n");
7     return NULL;
8 }
9
10 int main(int argc, char** argv)
11 {
12     pthread_t thread;
13
14     pthread_create(&thread, NULL, &hello, NULL);
15     printf("Hello World from main!\n");
16     pthread_join(thread, NULL);
17
18     return 0;
19 }
```

# Compile & Output

```
gcc --std=gnu99 -pthread -Wall  
    -o hello_world hello_world.c
```

```
Hello World from main!  
Hello World from pthread!
```

# More than One: Hello World with Pthreads Ver. 1

```
1 #include <stdlib.h>
2
3 ...
4
5 int main(int argc, char** argv)
6 {
7     int num_threads = 3; pthread_t *thread;
8     thread = malloc(num_threads * sizeof(*thread));
9
10    for (int i = 0; i < num_threads; i++)
11        pthread_create(thread + i, NULL, &hello, NULL );
12
13    for (int i = 0; i < num_threads; i++)
14        pthread_join(thread[i], NULL );
15
16    return 0;
17 }
```

# Output

```
[user]$ ./hello_world_2
Hello World from pthread!
Hello World from pthread!
Hello World from pthread!
```

# Single Argument: Hello World with Pthreads Ver. 2

```
1 void * hello(void *ptr)
2 {
3     int arg = *(int*)ptr;
4     printf("Hello World from pthread %d!\n", arg);
5     return NULL;
6 }
```

# Single Argument: Hello World with Pthreads Ver. 2

```
1 int main(int argc, char** argv)
2 {
3     int num_threads = 3;
4     pthread_t *thread;
5     int *arg;
6
7     thread = malloc(num_threads * sizeof(*thread));
8     arg = malloc(num_threads * sizeof(*arg));
9
10    for (int i = 0; i < num_threads; i++)
11    {
12        arg[i] = i;
13        pthread_create(thread + i, NULL, &hello, arg + i);
14    }
15
16    for (int i = 0; i < num_threads; i++)
17        pthread_join(thread[i], NULL);
18
19    free(thread);
20    free(arg);
21
22    return 0;
23 }
```

# Output

```
[user]$ ./hello_world_3  
Hello World from pthread 0!  
Hello World from pthread 1!  
Hello World from pthread 2!
```

# Many Arguments: Hello World with Pthreads Ver. 3

```
1 struct pthread_args
2 {
3     long thread_id;
4     long num_threads;
5 };
6
7 void * hello(void *ptr) {
8     struct pthread_args *arg = (struct pthread_args*) ptr;
9     printf("Hello World from %ld of %ld PID = %d TID = %li!\n",
10         arg->thread_id,
11         arg->num_threads,
12         getpid(),
13         pthread_self());
14
15     return NULL ;
16 }
```



# Many Arguments: Hello World with Pthreads Ver. 3

```
1 #include <unistd.h>
2 int main(int argc, char** argv)
3 {
4     long num_threads = 3;
5     pthread_t *thread;
6     struct pthread_args *arg;
7     thread = malloc(num_threads * sizeof(*thread));
8     arg = malloc(num_threads * sizeof(*arg));
9
10    for (int i = 0; i < num_threads; i++)
11    {
12        arg[i].thread_id = i;
13        arg[i].num_threads = num_threads;
14        pthread_create(thread + i, NULL, &hello, arg + i);
15    }
16
17    for (int i = 0; i < num_threads; i++)
18        pthread_join(thread[i], NULL);
19    free(thread);
20    free(arg);
21    return 0;
22 }
```

# Output

```
[user]$ ./hello_world_4  
Hello World from pthread 1 of 3 PID = 23750 TID = 23752!  
Hello World from pthread 0 of 3 PID = 23750 TID = 23751!  
Hello World from pthread 2 of 3 PID = 23750 TID = 23753!
```

# Return Result from Pthread in struct Argument

```
1 struct pthread_args
2 {
3     int in, out;
4 };
5
6 void * triple(void *ptr)
7 {
8     struct pthread_args *arg = ptr;
9     arg->out = 3 * arg->in;
10    return NULL ;
11 }
```

# Return Result from Pthread in struct Argument

```
1 int main(int argc, char** argv) {
2     int num_threads = 3; pthread_t *thread;
3     struct pthread_args *arg;
4
5     thread = malloc(num_threads * sizeof(*thread));
6     arg = malloc(num_threads * sizeof(*arg));
7
8     for (int i = 0; i < num_threads; i++){
9         arg[i].in = i;
10        pthread_create(thread + i, NULL, &triple, arg + i);
11    }
12
13    for (int i = 0; i < num_threads; i++){
14        pthread_join(thread[i], NULL);
15        printf("Triple of %d is %d\n",
16              arg[i].in,
17              arg[i].out);
18    }
19    free(thread);
20    free(arg);
21    return 0;
22 }
```

# Return Result from Pthread as Pointer to Memory

```
1
2 void * triple(void *ptr) {
3
4     int *out = malloc(sizeof(*out));
5     *out = 3 * (*(int*)ptr);
6
7     return (void*)out;
8 }
```

# Return Result from Pthread as Pointer to Memory

```
1 int main(int argc, char** argv) {
2     int num_threads = 3;
3     pthread_t *thread;
4     int *in;
5
6     thread = malloc(num_threads * sizeof(*thread));
7     in = malloc(num_threads * sizeof(*in));
8
9     for (int i = 0; i < num_threads; i++){
10         in[i] = i;
11         pthread_create(thread + i, NULL, &triple, in + i);
12     }
13
14     for (int i = 0; i < num_threads; i++){
15         int *out;
16         pthread_join(thread[i], (void*)&out);
17         printf("Triple of %d is %d\n", in[i], *out);
18         free(out);
19     }
20     free(thread);
21     free(in);
22     return 0;
23 }
```

# What have we covered so far?

- Creating new threads with `pthread_create`
- Waiting for threads to finish with `pthread_join`
- Passing arguments to a pthread function
- Returning results from pthread function

# Assignment: EMSim (EM Simulator)



## Task

- EMSim simulates the European Championship 2016 by utilizing results from the last 50 years.
- Every match will be played from the group- and final-phase.
- Your task is to parallelize the function `playEM()` and `playFinalRound()` so that the application has a speedup greater 2 with 4 threads.



# Assignment: EMSim - Usage

## Usage of the program

- Sequential:  
`./emsim_seq <database> 1`
- Parallel:  
`./emsim_par <database> (<#threads>)`

# Assignment: EMSim - playEM()

```
1 void playEM(team_t* teams) {
2     <... some declarations ...>
3
4     // play groups
5     initialize(); // necessary for the unit testing
6     for (g = 0; g < NUMGROUPS; ++g) {
7         playGroup(g, teams + (g * cTeamsPerGroup), cTeamsPerGroup,
8                 successors + g * 2, successors + (numSuccessors - (g * 2) - 1),
9                 bestThirds + g);
10    }
11
12    // fill best thirds
13    sortTeams(NUMGROUPS, bestThirds);
14    for (g = 0; g < numSuccessors; ++g)
15        if (successors[g] == NULL) successors[g] = bestThirds[curBestThird++];
16
17    // play final rounds
18    while (numSuccessors > 1) {
19        playFinalRound(numSuccessors / 2, successors, successors);
20        numSuccessors /= 2;
21    }
22 }
```

# Assignment: EMSim - playFinalRound()

```
1 void playFinalRound(int numGames, team_t** teams, team_t** successors) {
2     team_t* team1;
3     team_t* team2;
4     int i, goals1 = 0, goals2 = 0;
5
6     for (i = 0; i < numGames; ++i) {
7         team1 = teams[i*2];
8         team2 = teams[i*2+1];
9         playFinalMatch(numGames, i, team1, team2, &goals1, &goals2);
10
11         if (goals1 > goals2)
12             successors[i] = team1;
13         else if (goals1 < goals2)
14             successors[i] = team2;
15         else {
16             playPenalty(team1, team2, &goals1, &goals2);
17             if (goals1 > goals2)
18                 successors[i] = team1;
19             else
20                 nsuccessors[i] = team2;
21         }
22     }
23 }
```

# Assignment: EMSim - Provided Files

- Makefile
  - contains rules to build executables
  - available targets: parallel, sequential, all (default), clean
  - 'mode=debug make [target]' to build debug version, use 'make clean' before
- main.c
  - main function - argument handling + build teams + call playEM
- emsim.h
  - Header file for emsim.c and emsim\_\*.c
- emsim.c
  - Defines the simulator logic
- db.h / db.c
  - Header and definition for the database accesses
- emsim\_seq.c
  - Sequential version of playEM().
- student/emsim\_par.c
  - Implement the parallel version in this file

- em.db
  - Input data: The database containing all em results.
- libsqlite3.a
  - the slite3 library to read the database
  - there is also a libsqlite3\_32.a (in case you have a 32bit system) -> in that case, you have to modify the Makefile
- vis.h / vis.c
  - The visualization component
- unit<sub>test</sub>.c
  - The unit tests that execute both the serial and parallel version to compare results.

# Assignment: Extract, Build, and Run

1. Extract all files to the current directory  
`tar -xvf assignment1.tar.gz`
2. Build the program  
`make [sequential] [parallel] [unittest]`
  - sequential: build the sequential program
  - parallel: build the parallel program
  - unit<sub>t</sub>est : *buildstheunittests*
3. Run the sequential program (100 repetitions)  
`student/emsim_seq em.db 1`
4. Run the parallel program (with N threads)  
`student/emsim_par em.db N`

# Submission

1. Log into the website
2. Go to Assignments
3. Use link for Assignment 1
4. Upload your `emsim_par.c`
5. Press Submit

Parallel Programming
Home
Slides
**Assignments**
Contact

Welcome ga49qez [Logout](#)

## NOTE:

1. Some of the gcc versions does not support -Wpedantic compiler flag. If you get an error saying Unrecognized command line option "-Wpedantic", change the flag to -pedantic in Makefile
2. Add -lrt option to LDFLAGS in case you get Undefined reference to "clock\_gettime" error.

 no file selected

Assignment Name	Submitted By	Date Of Submission	Status
Histogram Pthreads	Andreas Wilhelm	April 20, 2015, 6:20 p.m.	✓ Accepted

```

Build step successfull
Correctness checks successfull
Pthread checks successfull
Memcheck checks successfull
Helgrind checks successfull
Continous checks successfull
Speedup: 4.0

```