


Parallel Programming Tutorial - Pthread 3

M.Sc. Andreas Wilhelm

Technical University Munich

May 02, 2016



TUM Uhrenturm

Organization

- There is no tutorial next week and the week after (May, 09 and 16)
- Deadline for Assignment 2 is on May 23, 3:30 pm
- The assignment will be online on Wednesday, May 04
- The slides will be updated then

Assignment 1: Possible Solution (Speedup 2,38) (1/3)

1. Parallelization of playEM()

- playEM() creates numThreads threads
- The calls of playGroup() are evenly distributed throughout the threads
- The arrays teams, successors, and bestThirds are given as input to each thread
- There is no race condition since each thread is operating on disjunct data
- bestThirds are sorted after the threads have been joined

```
1 struct pthread_args
2 {
3     int groupMin; // start index of the group number
4     int groupMax; // end index of the group number
5     int numSuccessors;
6     int teamsPerGroup;
7     team_t* teams;
8     team_t** successors;
9     team_t** bestThirds;
10 };
```

Assignment 1: Possible Solution (Speedup 2,38) (2/3)

```

1 // compute number of groups for each thread (plus a rest)
2 int g = 0
3 int groupsPerThread = NUMGROUPS / numThreads;
4 int groupsRest = NUMGROUPS % numThreads;
5
6 // assign the groups to numThreads threads
7 for (t = 0; t < numThreads; ++t) {
8     thread_arg[t].groupMin = g;
9     g += groupsPerThread;
10    if (groupsRest > 0) { g++; groupsRest--; }
11    thread_arg[t].groupMax = g;
12    thread_arg[t].teams = teams;
13    thread_arg[t].numSuccessors = numSuccessors;
14    thread_arg[t].teamsPerGroup = teamsPerGroup;
15    thread_arg[t].successors = successors;
16    thread_arg[t].bestThirds = bestThirds;
17    pthread_create(&threads[t], NULL, &parallel_calls_to_playGroup, thread_arg+t);
18 }

1 void* parallel_calls_to_playGroup(void * ptr) {
2     struct pthread_args *arg = ptr;
3     int g;
4
5     for (g = arg->groupMin; g < arg->groupMax; ++g) {
6         playGroup(g, arg->teams + (g * arg->teamsPerGroup),
7                 arg->teamsPerGroup,
8                 arg->successors + g * 2,
9                 arg->successors + (arg->numSuccessors - (g * 2) - 1),
10                arg->bestThirds + g);
11    }
12    return NULL;
13 }

```

Assignment 1: Possible Solution (Speedup 2,38) (3/3)

2. Parallelization of playFinalRound()

- Each call of `playFinalRound()` creates `numSuccessors` threads
- Each thread calls `playFinalMatch()` for a single match
- Input is the number of games, the current game number, and the team information
- The resulting number of goals is written to the pointer argument by each thread

```


1 struct pthread_args_match {
2     int numGames;
3     int gameNo;
4     team_t* team1;
5     team_t* team2;
6     int goals1;
7     int goals2;
8 };

1 for (i = 0; i < numGames; ++i) {
2     args[i].numGames = numGames;
3     args[i].gameNo = i;
4     args[i].team1 = teams[i*2];
5     args[i].team2 = teams[i*2+1];
6     pthread_create(&threads + i, NULL, &playMatchInPar, &args[i]);
7 }

1 void* playMatchInPar(void* ptr) {
2     struct pthread_args_match *args = (ptr);
3     playFinalMatch(args->numGames, args->gameNo,
4                   args->team1, args->team2, &args->goals1, &args->goals2);
5     return NULL;
6 }

```

Assignment 1: Fast solution (Speedup 3,41) (1/2)



1. Parallelization of playEM()

- Idea: avoid load imbalance by a clever distribution of the `playGroup()` calls
- Thread 1: `playGroup(0)`
- Thread 2: `playGroup(1, 3)`
- Thread 3: `playGroup(2, 5)`
- Thread 4: `playGroup(4)`

Assignment 1: Fast solution (Speedup 3,41) (2/2)

2. Parallelization of playFinalMatch() calls

- The results for all games (group and final games) are deterministic
- Idea: play the final matches in parallel to the group games by additional 4 threads
- Postpone the join of all threads to the end of playEM()

```

1 void* parallel_calls_to_playFinalMatch(void* arg){
2     struct pthread_args_final* args = arg;
3     int goals1, goals2, i;
4
5     for (i = 0; i < args->threadGames; ++i) {
6         playFinalMatch(args->numGames, args->gameNo,
7                         args->team1, args->team2,
8                         &goals1, &goals2);
9         if (goals1 == goals2)
10             playPenalty(args->team1, args->team2, &goals1, &goals2);
11         args++;
12     }
13     return NULL;
14 }

```

```

1 finalArgs[0].threadGames = 4;
2 finalArgs[0].numGames = 8;
3 finalArgs[0].gameNo = 0;
4 finalArgs[0].team1 = teams + 0;
5 finalArgs[0].team2 = teams + 1;
6 // ...
7 finalArgs[12].threadGames = 3;
8 finalArgs[12].numGames = 2;
9 finalArgs[12].gameNo = 0;
10 finalArgs[12].team1 = teams + 0;
11 finalArgs[12].team2 = teams + 8;

```

Condition Variables

Definition

A Condition Variable (CV) is the mechanism your program waits for a predicate to become true, and to communicate to others that it might be true

- Used for communicating information about shared state data
- Some threads may **wait** for some predicate to become true
- Some threads may **signal** a single waiting thread
- Some threads may **broadcast** all waiting threads
- Example: Signal that a queue is no longer empty
- Every condition variable is associated with:
 - a mutex (for synchronization)
 - a predicate (the shared state data)

Condition Variables: Initializing

```
1 pthread_cond_t cond = PTHREAD_COND_INITIALIZER;  
2 int pthread_cond_init ( pthread_cond_t *cond,  
3                       pthread_condattr_t *condattr );  
4 int pthread_cond_destroy( pthread_cond_t *cond );
```

- `pthread_cond_t *cond`
 - Pointer to a condition variable
- `pthread_condattr_t *condattr`
 - Optional pointer to a `pthread_condattr_t` struct to define behaviour, if NULL defaults are used
- Use static initialization for static CV's with default attributes (you do not have to destroy a static CV)
- Use dynamic initialization for malloc and non-standard attributes (destroying of the CV necessary)

Condition Variables: Waiting (1/2)

```
1 int pthread_cond_wait( pthread_cond_t *cond,  
2                       pthread_mutex_t *mutex );  
3 int pthread_cond_timedwait( pthread_cond_t *cond,  
4                             pthread_mutex_t *mutex,  
5                             struct timespec *expiration );
```

- `pthread_mutex_t *mutex`
 - Pointer to the mutex that is associated with cond
- `struct timespec *expiration`
 - The timespec structure that sets the max. expiration time
- Before waiting on a CV, the mutex must always be locked
- Waiting on a CV unlocks the mutex and waits on a signal/broadcast (atomically!)
- On signal/broadcast, the mutex will be locked again

Condition Variables: Waiting (2/2)

```
1 int pthread_cond_wait( pthread_cond_t *cond,  
2                       pthread_mutex_t *mutex );  
3 int pthread_cond_timedwait( pthread_cond_t *cond,  
4                             pthread_mutex_t *mutex,  
5                             struct timespec *expiration );
```

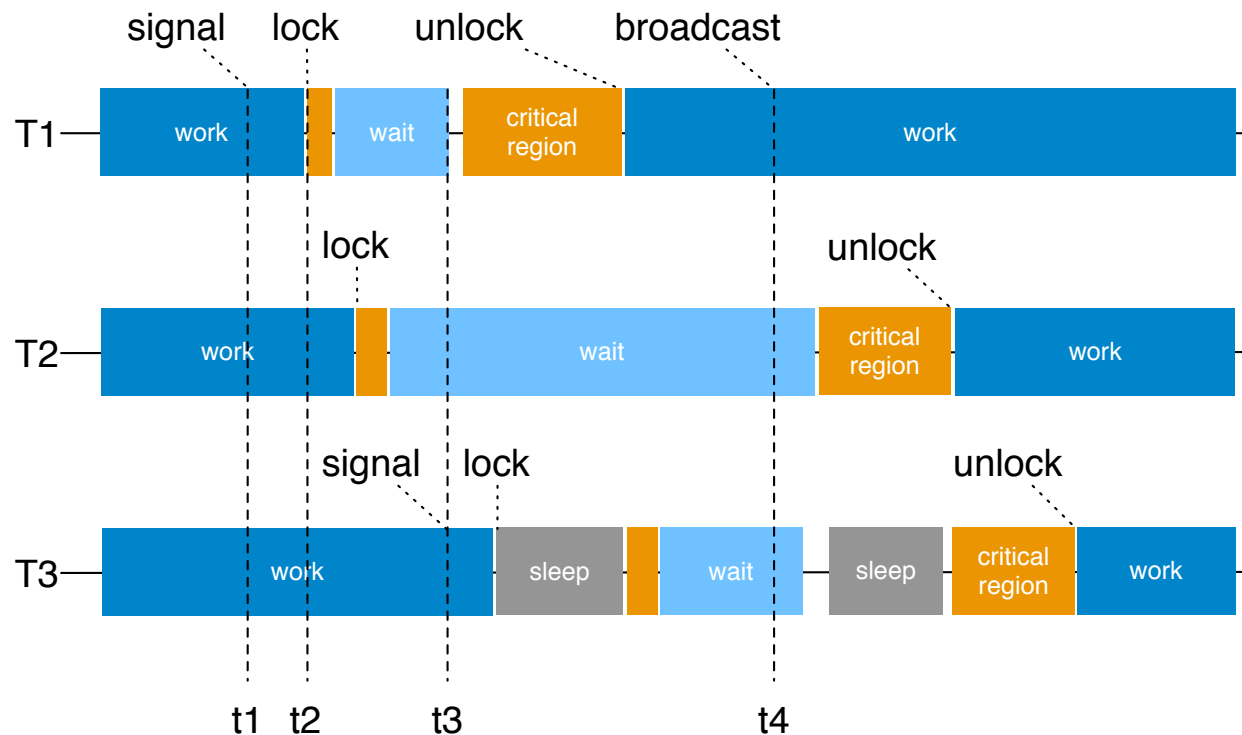
- Always check your predicate before waiting and then check it again (in a loop) to prevent:
 1. Intercepted Wakeups
Other threads may wake up earlier and grab the work. Never assume that the last `pthread_cond_wait` gets the lock!
 2. Loose Predicates
Signaling thread may use a weak predicate ("there may be work" instead of "there is work"). You have to check for strict predicate again!
 3. Spurious Wakeups
Wait may wakeup without signal/broadcast!

Condition Variables: Waking up Waiters (1/2)

```
1 int pthread_cond_signal( pthread_cond_t *cond );  
2 int pthread_cond_broadcast( pthread_cond_t *cond );
```

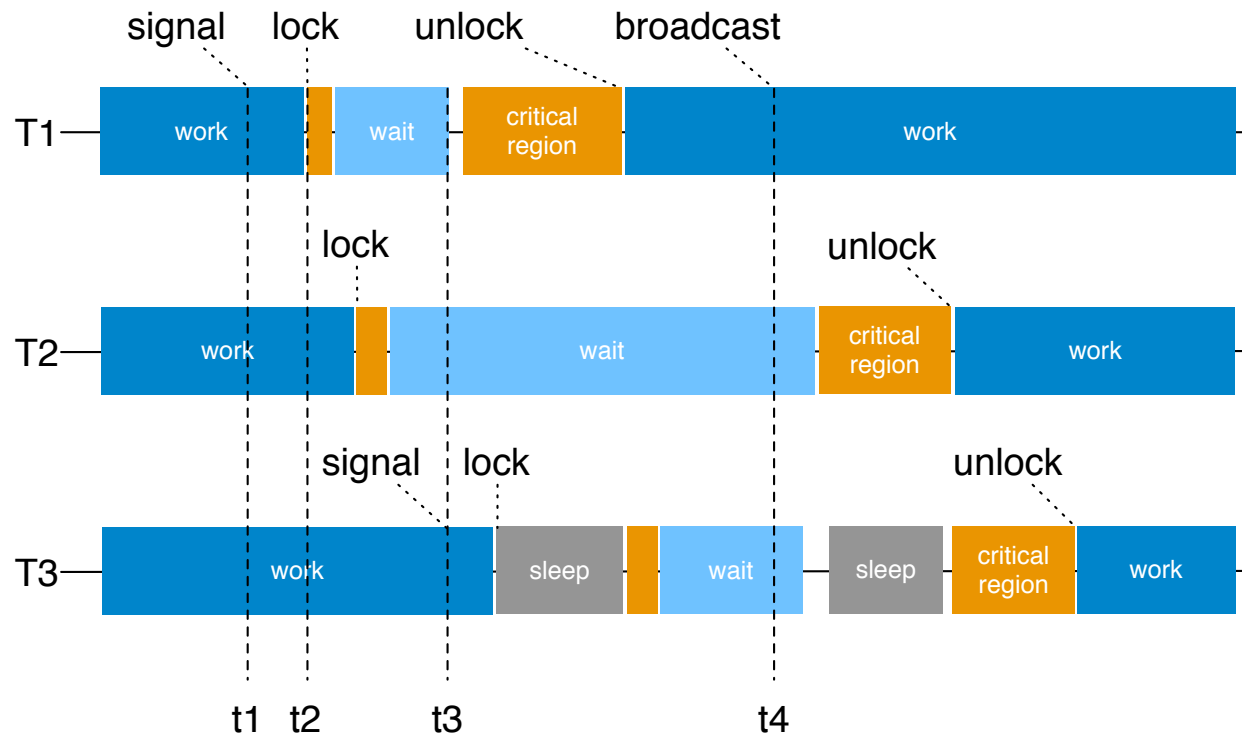
- signal wakes up a single thread waiting on cond
- broadcast wakes up all threads waiting on cond
- signal/broadcast without any waiters has no effect
- signal is more efficient (less context switches)
- signal/broadcast while holding the mutex may infer additional context switches, but may prevent intercepted wakeups

Condition Variables: Illustration




- **t_1 :** T1 signals with no waiters

Condition Variables: Illustration




- **t2:** T1 waits for a predicate to hold. Therefore, it first locks a mutex, then checks the predicate (false), and then calls `pthread_cond_wait` (which unlocks the mutex and waits)

Condition Variables: Illustration



- **t3:** T3 signals and wakes up T1, which locks the mutex and works on the critical region.

Condition Variables: Illustration



- **t4:** T1 broadcasts and wakes up T2 and T3.

Condition Variables: Example (1/2)

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include <time.h>
4  #include <errno.h>
5
6  typedef struct {
7      pthread_mutex_t  mutex;
8      pthread_cond_t   cond;
9      int              value;
10 } my_struct_t;
11
12 my_struct_t data = { PTHREAD_MUTEX_INITIALIZER,
13                     PTHREAD_COND_INITIALIZER,
14                     0 };
15
16 int hibernation = 1; // Default to 1 second
17
18 void *signal_thread( void *arg ) {
19
20     sleep( hibernation );
21     pthread_mutex_lock( &data.mutex );
22     data.value = 1; // set predicate
23     pthread_cond_signal( &data.cond );
24     pthread_mutex_unlock( &data.mutex );
25
26     return NULL;
27 }
```

Condition Variables: Example (2/2)

```
1 int main( int argc, char**argv ) {
2
3     pthread_t thread;
4     struct timespec timeout;
5
6     pthread_create( &thread, NULL, &signal_thread, NULL);
7     clock_gettime(CLOCK_REALTIME, &timeout);
8     timeout.tv_sec += 2;    // wait for predicate for 2 seconds
9
10    pthread_mutex_lock( &data.mutex );
11    while( data.value == 0) { // important!
12        if (pthread_cond_timedwait( &data.cond, &data.mutex, &timeout ) == ETIMEDOUT) {
13            printf( "Condition wait timed out.\n");
14            break;
15        }
16    }
17    if (data.value != 0) // consider the timeout
18        printf( "Condition was signaled.\n");
19    pthread_mutex_unlock( &data.mutex );
20
21    pthread_join( thread, NULL );
22    return 0;
23 }
```

Pthreads: What is missing?

- Special attributes for threads, mutexes, etc.
- Thread cancelation
- Thread-specific data
- Scheduling
- Priority-aware mutexes
- Barriers, Semaphores, Read-Write-Locks, etc.

Assignment 2: Dynamic EMSim (1/2)

Task

- Use POSIX threads to parallelize `playGroups()` and `playFinalRound()` of `emsim_seq.c`
- The program should follow the producer/consumer pattern:
 - There is one producer thread (the main thread) that produces game information and writes it to a buffer (for group games and final games).
 - There are `numWorker` worker threads that use the game information as input and play the matches asynchronously.
 - As soon as match information is available on the buffer, a free worker grabs it and begins to play.
 - The results are used by the producer thread to progress.
- Consider:
 - This time, `playGroups()` will be called once for all 36 matches of the group phase.
 - The number of created worker threads `N` is checked (`numWorker <= N <= 2 x numWorker`).
 - You may have to use global storage, think about synchronization.
 - You can solve the assignment with or without condition variables.
 - The speedup with 4 cores must be at least 3.
 - There is a new executable (`checks`) that you can use for valgrind/ltrace checks (use debug mode to get meaningful messages).

Assignment 2: Dynamic EMSim (2/2)

Hints for using condition variables

- Use two condition variables:
 - One CV for the producer (consumer ready)
 - One CV for the consumer (work available)
- Use a single mutex for both condition variables.
- The producer needs to wait for the consumers to finish their work at the end.
- Remember: If there is a signal and no one is waiting -> nothing happens.
- Never make a copy of a CV, only references allowed.

Assignment 2: emsim_seq.c (1/2)

```

1  // play all (36) group games by utilizing numWorker worker threads
2  void playGroups(team_t* teams, int numWorker) {
3      static const int cNumTeamsPerGroup = NUMTEAMS / NUMGROUPS;
4      int g, i, j, goalsI, goalsJ;
5
6      for (g = 0; g < NUMGROUPS; ++g) {
7          for (i = g * cNumTeamsPerGroup; i < (g+1) * cNumTeamsPerGroup; ++i) {
8              for (j = (g+1) * cNumTeamsPerGroup - 1; j > i; --j) {
9
10                 // team i plays against team j in group g
11                 playGroupMatch(g, teams + i, teams + j, &goalsI, &goalsJ);
12                 teams[i].goals += goalsI - goalsJ;
13                 teams[j].goals += goalsJ - goalsI;
14                 if (goalsI > goalsJ)
15                     teams[i].points += 3;
16                 else if (goalsI < goalsJ)
17                     teams[j].points += 3;
18                 else {
19                     teams[i].points += 1;
20                     teams[j].points += 1;
21                 }
22             }
23         }
24     }
25 }

```

Assignment 2: emsim_seq.c (2/2)

```
1 // play a specific final round by utilizing numWorker worker threads
2 void playFinalRound(int numGames, team_t** teams, team_t** successors, int numWorker){
3     team_t* team1;
4     team_t* team2;
5     int i, goals1 = 0, goals2 = 0;
6
7     for (i = 0; i < numGames; ++i) {
8         team1 = teams[i*2];
9         team2 = teams[i*2+1];
10        playFinalMatch(numGames, i, team1, team2, &goals1, &goals2);
11
12        if (goals1 > goals2)
13            successors[i] = team1;
14        else if (goals1 < goals2)
15            successors[i] = team2;
16        else {
17            playPenalty(team1, team2, &goals1, &goals2);
18            if (goals1 > goals2)
19                successors[i] = team1;
20            else
21                successors[i] = team2;
22        }
23    }
24 }
```

Assignment: Dynamic EMSim - Provided Files

- Makefile
 - contains rules to build executables
 - available targets: parallel, sequential, unit_test, checks, all (default), clean
 - 'mode=debug make [target]' to build debug version, use 'make clean' before
- main.c
 - main function - argument handling + build teams + call playEM
- emsim.h
 - Header file for emsim.c and emsim_*.c
- emsim.c
 - Defines the simulator logic
- db.h / db.c
 - Header and definition for the database accesses
- emsim_seq.c
 - Sequential version of playGroups() and playFinalRound().
- student/emsim_par.c
 - Implement the parallel version in this file

Assignment: Dynamic EMSim - Provided Files

- em.db
 - Input data: The database containing all em results.
- libsqlite3.a
 - the slite3 library to read the database
 - there is also a libsqlite3_32.a (in case you have a 32bit system) -> in that case, you have to modify the Makefile
- vis.h / vis.c
 - The visualization component
- unit_test.c
 - The unit tests that execute both the serial and parallel version to compare results.