

Parallel Programming Tutorial - OpenMP 1

M.Sc. Andreas Wilhelm
 Technical University Munich
 May 23, 2016



TUM Uhrenturm

Organization

Schedule

- May 23: OpenMP 1
- May 30: OpenMP 2
- June 6: Dependence Analysis
- June 13: Loop Transformations
- June 20: MPI 1
- June 27: MPI 2
- July 4: Question session

Assignment 2: Possible Solution (Speedup 3) (1/7)

1. Parallelize playGroups()

- use the main thread as producer (playGroups())
- utilize N worker threads which play both group and final games
- use a global buffer of fixed size to communicate with worker threads
- use a global index that points to the next unprocessed buffer entry
- use a global counter for the number of processed matches
- the buffer contains a mode flag to distinguish group and final games

```
1 #define BUFFER_SIZE 50
2
3 typedef struct {
4     int mode;    // 0 for group game, 1 for final game
5     int number_games;
6     int group_number;
7     team_t *team1;
8     team_t *team2;
9     team_t **successors;
10 } match_info;
11
12 match_info *match_buffer;
13 int buffer_index = -1;
14 int completed_games = 0;
```

Assignment 2: Possible Solution (Speedup 3) (2/7)

1. Parallelize playGroups()

- wait for the consumer (condition variable `need_work`) if the buffer is full
- fill the buffer with match information and signal that new work (condition variable (`new_work`) is available

```

16 void playGroups(team_t* teams, int numWorker) {
17     //... some declarations
18     match_buffer = malloc(BUFFER_SIZE * sizeof(match_info));
19
20     for (g = 0; g < numWorker; ++g)
21         pthread_create(threads + g, NULL, &consumer_game, NULL);
22
23     for (g = 0; g < NUMGROUPS; ++g)
24         for (i = g * cNumTeamsPerGroup; i < (g+1) * cNumTeamsPerGroup; ++i)
25             for (j = (g+1) * cNumTeamsPerGroup - 1; j > i; --j) {
26                 pthread_mutex_lock(&mutex);
27                 while (buffer_index >= BUFFER_SIZE) {
28                     pthread_cond_wait(&need_work, &mutex);
29                 }
30
31                 match_buffer[++buffer_index].group_number = g;
32                 match_buffer[buffer_index].team1 = teams + i;
33                 match_buffer[buffer_index].team2 = teams + j;
34                 match_buffer[buffer_index].mode = 0;
35
36                 pthread_cond_signal(&new_work);
37                 pthread_mutex_unlock(&mutex);
38             }
39     //... see next slide

```

Assignment 2: Possible Solution (Speedup 3) (3/7)

1. Parallelize playGroups()

- wait for the consumers to finish work after producing all matches
- broadcast to all consumers `new_work`
- wait for consumers to increment `completed_games` so that it is equal to `numWorker + 1`

```
40 pthread_mutex_lock(&mutex);
41 completed_games = 1;
42 pthread_cond_broadcast(&new_work);
43
44 for (i = 0; i < numWorker; ++i) {
45     if (completed_games == numWorker+1) {
46         break;
47     }
48     pthread_cond_wait(&work_done, &mutex);
49 }
50 pthread_mutex_unlock(&mutex);
51 }
```

Assignment 2: Possible Solution (Speedup 3) (4/7)

2. Parallelize playFinalRound()

- terminate all consumer threads in case of the final match (send broadcast)
- play final match
- join threads

```
40 void playFinalRound(int numGames, team_t** teams, team_t** successors, int numWorker) {
41     pthread_mutex_lock(&mutex);
42     completed_games = 0;
43     pthread_mutex_unlock(&mutex);
44
45     if (numGames == 1) {
46         pthread_mutex_lock(&mutex);    // stop remaining threads
47         terminate = 1;
48         pthread_cond_broadcast(&new_work);
49         pthread_mutex_unlock(&mutex);
50
51     // <play final match>
52
53     for (i = 0; i < numWorker; ++i) {
54         pthread_join(threads[i], NULL);
55     }
56 }
57 // continue on next slides
```

Assignment 2: Possible Solution (Speedup 3) (5/7)

2. Parallelize playFinalRound()

— for all other finals: wait for consumers, fill buffers and send new_work

```

40 void playFinalRound(int numGames, team_t** teams, team_t** successors, int numWorker) {
41 // ...
42 else {
43     for (int i = 0; i < numGames; ++i) {
44         pthread_mutex_lock(&mutex);
45         while (buffer_index >= BUFFER_SIZE) {
46             pthread_cond_wait(&need_work, &mutex);
47         }
48
49         match_buffer[++buffer_index].number_games = numGames;
50         match_buffer[buffer_index].mode = 1;
51         match_buffer[buffer_index].group_number = i;
52         match_buffer[buffer_index].team1 = teams[2*i];
53         match_buffer[buffer_index].team2 = teams[2*i+1];
54         match_buffer[buffer_index].successors = successors;
55
56         pthread_cond_signal(&new_work);
57         pthread_mutex_unlock(&mutex);
58     }
59 // ... continue on next slide

```

Assignment 2: Possible Solution (Speedup 3) (6/7)

2. Parallelize playFinalRound()

— wait for completed games to be finished

```
40 void playFinalRound(int numGames, team_t** teams, team_t** successors, int numWorker) {
41     // ...
42     pthread_mutex_lock(&mutex);
43     completed_games = 1;
44     pthread_cond_broadcast(&new_work);
45     for (i = 0; i < numWorker; ++i) {
46         if (completed_games == numWorker+1) {
47             break;
48         }
49         pthread_cond_wait(&work_done, &mutex);
50     }
51     pthread_mutex_unlock(&mutex);
52 }
53 }
```


Assignment 2: Possible Solution (Speedup 3) (7/7)

3. consumer_game()

— use a global flag to terminate consumer thread (int terminate)

```
40 int terminate = 0;
41
42 void *consumer_game(void *arg) {
43     //...
44     while (1) {
45         pthread_mutex_lock(&mutex);
46
47         while (buffer_index < 0) {
48             if (terminate) { // check if consumer sent a signal to terminate
49                 pthread_mutex_unlock(&mutex);
50                 return NULL;
51             }
52
53             if (completed_games) { // increment completed games for
54                 ++completed_games;
55                 pthread_cond_signal(&work_done);
56             }
57             pthread_cond_wait(&new_work, &mutex);
58         }
59
60         info = match_buffer[buffer_index--]; // get next match to play
61         pthread_cond_signal(&need_work); // signal that consumer consumed a match
62         pthread_mutex_unlock(&mutex);
63
64         if (info.mode) // play final match
65             else // play group match
66     }
```

Introduction to OpenMP

- OpenMP is an API for explicit shared-memory parallelism
- Supported by most compilers (gcc, icc, msvc, clang)
- Utilizes OS threading capabilities (e.g. Pthreads)
- Fully documented in the specification (see <http://www.openmp.org/mp-documents/OpenMP4.0.0.pdf>)
- Comprised of three programming layer components

1. Compiler Directives

- Spawning parallel regions
- Distributing loop iterations across threads
- Synchronization
- ...

2. Runtime Library Routines

- Setting/Querying the number of current threads
- Querying thread-id's and wall-clock time
- ...

3. Environment Variables

- Setting number of threads
- Binding threads to processors
- ...

Directives

Format

`#pragma omp <directive name> <{clause, ...}>`

- `#pragma omp`
Required for all OpenMP C/C++ directives
- `directive name`
A valid OpenMP directive
- `{clause, ...}`
Optional. Clauses can be in any order
- Most OpenMP constructs apply to a structured block


Example

```
#pragma omp parallel default(shared) private(i)
```

Parallel Region

```
#pragma omp parallel <{clause, ...}>
```

- A block of code that will be executed by multiple threads
- Number of threads defined by `#cpu`, clauses, or env. variables
- The reaching thread (0) creates a team of N threads (1, ..., N-1)
- At the end of a block there is an implicit join (barrier)
- There may be nested parallel regions



Parallel Region - Example

- `omp_get_thread_num()` returns the current thread number
- `omp_get_num_threads()` returns the number of threads

```
1 #include <omp.h>
2 #include <stdio.h>
3
4 int main(int argc, char** argv) {
5
6     #pragma omp parallel
7     {
8         printf("Hello World from thread %d\n", omp_get_thread_num());
9
10        // only executed by main thread
11        if (omp_get_thread_num() == 0)
12            printf("Number of threads is %d\n", omp_get_num_threads());
13    }
14    return 0;
15 }
```

`./hello_world`

Hello World from thread 1

Hello World from thread 0

Number of threads is 3

Hello World from thread 2

Parallel Region - Clauses

- `if (<scalar expression>)`
only executed multithreaded if scalar expr. evaluates to non-zero
- `private (<list>)`
each thread gets a copy of variables in a comma separated list (variables might be uninitialized)
- `firstprivate/lastprivate (<list>)`
same as `private`, but value is copied at the entry/exit
- `shared (<list>)`
variables in `list` are shared (no elements of structs or arrays)
- `default (shared | none)`
sets the default behaviour (`none` means that data sharing needs to be explicit)
- `reduction (<operator: list>)`
reduction operation and associated operand
- `num_threads (<integer expression>)`
sets the number of threads for the parallel region
- ...

for Directive

```
#pragma omp for <{clause, ...}>
```

- Worksharing construct to execute the immediately following loop by a team of threads
- Assumes that a parallel region has already been initiated
- There is an implicit barrier at the end of the loop
- Clauses:
 - `schedule (static|dynamic|guided|runtime|auto)`
sets the scheduling behaviour (see next slide)
 - `nowait`
threads do not synchronize after the parallel loop
 - `ordered`
iterations must have the same order as in a serial program
 - `collapse`
specifies the number of (nested) loops that shall be collapsed into a larger iteration space
 - `private, firstprivate...`

schedule clause

- `schedule (static, chunk_size)`

The iterations are divided into chunks of size *chunk_size* and assigned to the threads in round-robin fashion. When no chunk size is specified, the iterations are equally divided (at most one iteration per thread).

- `schedule (dynamic, chunk_size)`

The iterations are distributed to threads in chunks as the executing threads request them. Each thread executes a chunk of iterations, then requests another chunk, until no chunks remain. Each chunk contains *chunk_size* iterations, except for the last chunk. If no chunk size is specified, it defaults to 1.

- `schedule (guided, chunk_size)`

Similar to `dynamic`, but...

At the beginning the size of each chunk is proportional to the number of unassigned iterations divided by the number of threads, decreasing to 1 (or *chunk_size*, if specified).

- `schedule (auto)`

The scheduling decision is given to the compiler/runtime system.

- `schedule (runtime)`

The scheduling decision is deferred until run time, the schedule and chunk size are taken from internal control variables.

for Directive - Example (1/5)

```

1  #include <omp.h>
2  #include <stdio.h>
3  #include <string.h>
4  #include <stdlib.h>
5  #include <unistd.h>
6
7  const char *colored_digit[] = {
8      "\e[1;30;1m0", "\e[1;31;1m1", "\e[1;32;1m2", "\e[1;33;1m3", "\e[1;34;1m4", "\e[1;35;1m5", "\e[1;36;1m6", "\e[1;37;1m7"
9  };
10
11 int main(int argc, char** argv) {
12     unsigned int x_size = 80;
13     unsigned int y_size = 40;
14     unsigned long str_len = strlen (colored_digit [0]);
15     char *string_2D = (char*)malloc(x_size * y_size * str_len + y_size);
16
17     #pragma omp parallel for schedule(runtime)
18     for (unsigned long i = 0; i < y_size; i++) {
19         for (unsigned int j = 0; j < x_size; j++) {
20             memcpy(string_2D + ( i * x_size * str_len + i ) + (j * str_len), colored_digit[omp_get_thread_num()], str_len );
21         }
22     }
23
24     for (unsigned int i = 0; i < y_size; i++) {
25         unsigned long row = i * x_size * str_len + i ;
26         printf("%s\n", string_2D + row);
27     }
28
29     printf("\033[0m");
30     return 0;
31 }

```

for Directive - Example (2/5)

```
OMP_NUM_THREADS=4 OMP_SCHEDULE="STATIC" ./scheduling
```

[illegible]

for Directive - Example (3/5)

```
OMP_NUM_THREADS=4 OMP_SCHEDULE="STATIC,2" ./scheduling
```

[illegible]

for Directive - Example (4/5)

```
OMP_NUM_THREADS=4 OMP_SCHEDULE="DYNAMIC" ./scheduling
```

[illegible]

for Directive - Example (5/5)

```
OMP_NUM_THREADS=4 OMP_SCHEDULE="GUIDED" ./scheduling
```

[illegible]

Assignment 3: EMSim with OpenMP

Task

- Use OpenMP to parallelize the loops in `playGroups()` and `playFinalRound()` of `emsim_seq.c`
- Your solution should have a speedup greater 2.5 on 4 cores
- Consider:
 - Avoid race conditions
 - Pick proper scheduling strategy

Assignment 3: EMSim with OpenMP - Provided Files

- Makefile
 - contains rules to build executables
 - available targets: parallel, sequential, unit_test, checks, all (default), clean
 - 'mode=debug make [target]' to build debug version, use 'make clean' before
- main.c
 - main function - argument handling + build teams + call playEM
- emsim.h
 - Header file for emsim.c and emsim_*.c
- emsim.c
 - Defines the simulator logic
- db.h / db.c
 - Header and definition for the database accesses
- emsim_seq.c
 - Sequential version of playGroups() and playFinalRound().
- student/emsim_par.c
 - Implement the parallel version in this file

Assignment 3: EMSim with OpenMP - Provided Files

- em.db
 - Input data: The database containing all em results.
- libsqlite3.a
 - the slite3 library to read the database
 - there is also a libsqlite3_32.a (in case you have a 32bit system) -> in that case, you have to modify the Makefile
- vis.h / vis.c
 - The visualization component
- unit_test.c
 - The unit tests that execute both the serial and parallel version to compare results.